

Rebuilding the **Crossword** with Google Cloud

O'Reilly Software Architecture Conference - NYC 2018

JP Robinson

whoami

Principal Software Engineer - Games

- NYT since 2011
- Using & evangelizing Go since 2013
 - “Using Go at The New York Times” - 2015
 - OSS: Gizmo, openapi2proto
- Google Cloud Platform
 - OSS: Marvin, drone-gae, spannerr
- Member, NYT Architecture Review Board

The NYT Crossword

- Started in 1942, daily in 1950
- Went digital in 1996, in-house in 2014
- Now available on Android, iOS, Web, Windows
- Daily (Free!) Mini crossword
- 300,000+ paid subscribers
- New puzzles are published nightly


The NYT Crossword Spike

Last 7 Days of Traffic



The NYT Crossword Spike

Last 7 Days of Traffic

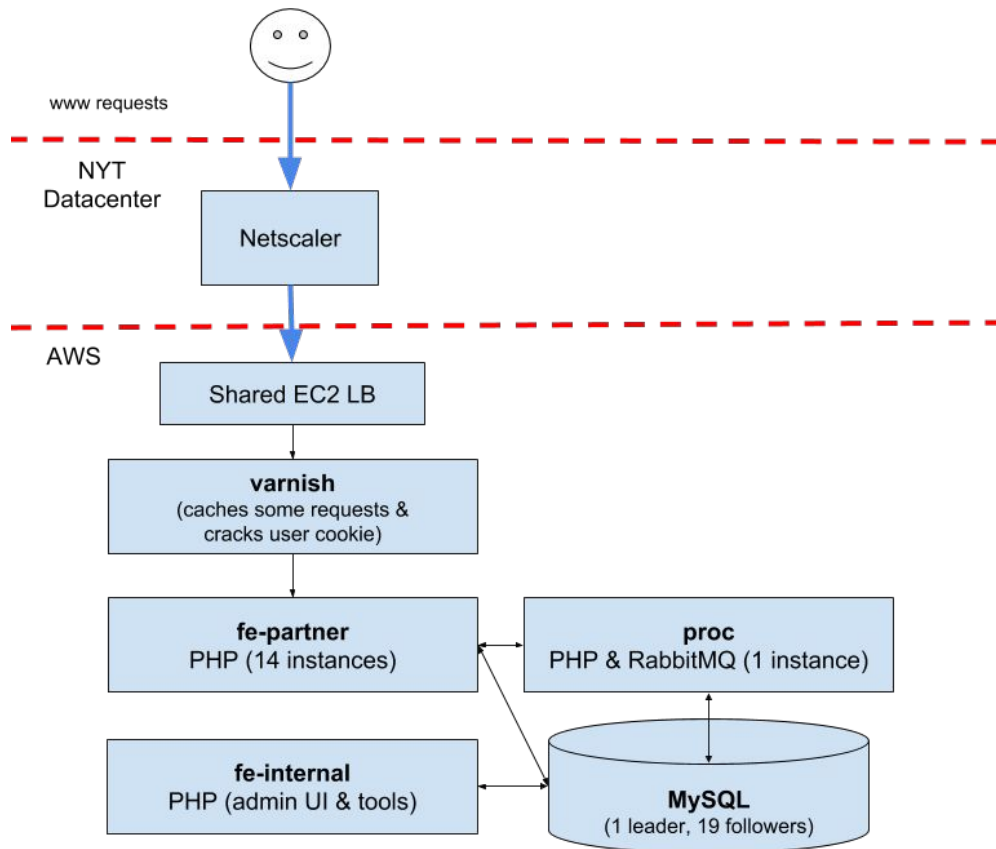


NYT's "Legacy" Infrastructure

Pre-2016:

- Multiple data centers
 - Core news product and internal services
- AWS
 - Search, Blogs, Interactive Articles, Cooking, Games
 - Internal console (Nimbul) for accessing services
- Mix of Puppet, Ansible and custom tooling

The “Legacy” Architecture




The “Legacy” Architecture

- Always scaled for peak traffic
- Adding VMs a manual process
- Deployments take 30-60+ minutes
- No log aggregation, minimal monitoring
- Internal PHP framework
- Several SPOF

Deciding on Go

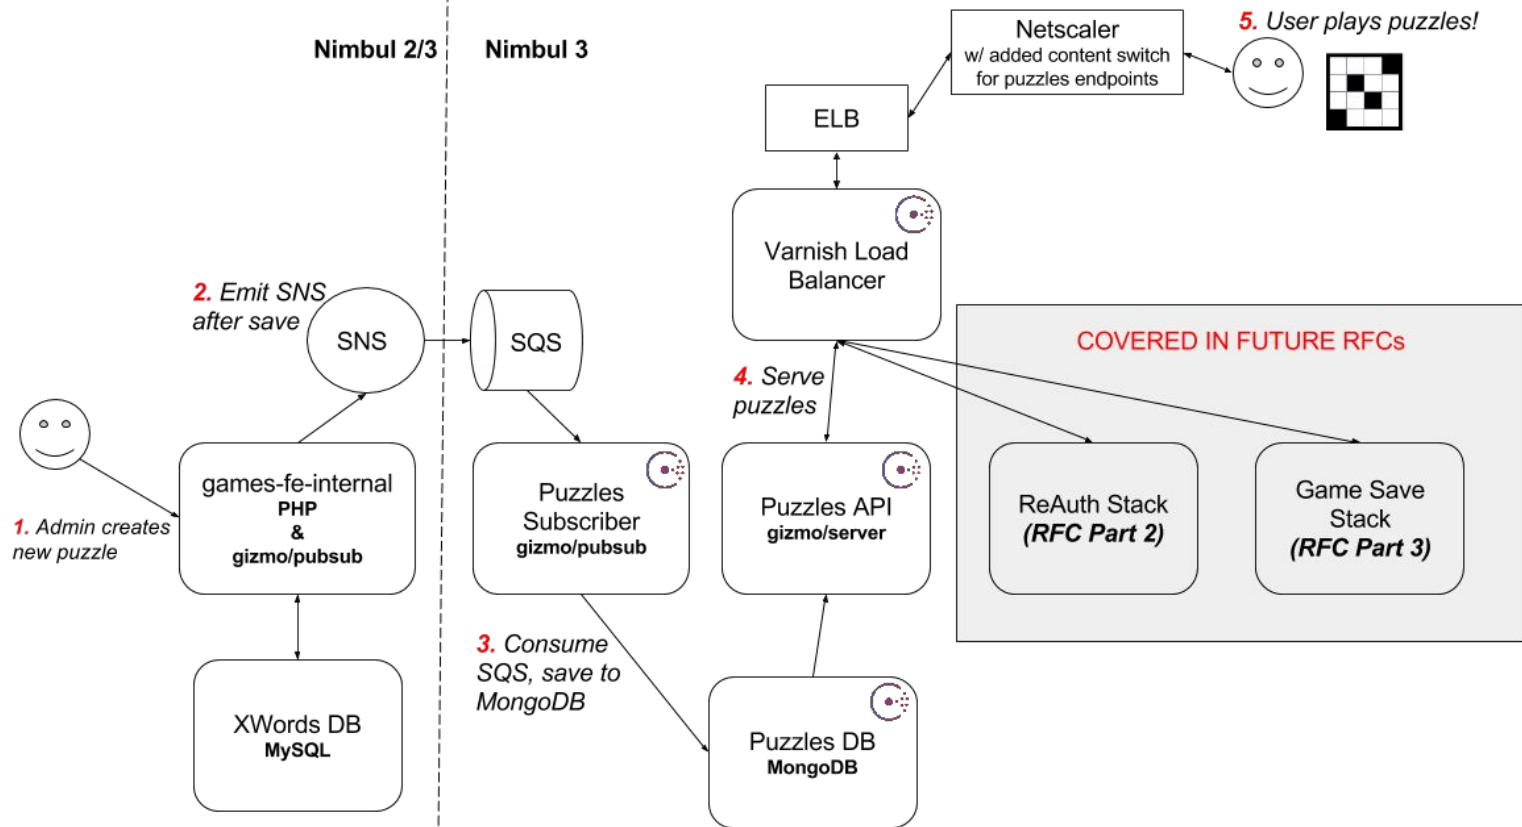
- Easy to read
- Easy to test
- Easy to deploy
- Efficient, fast and fun
- 9+ NYT teams using it



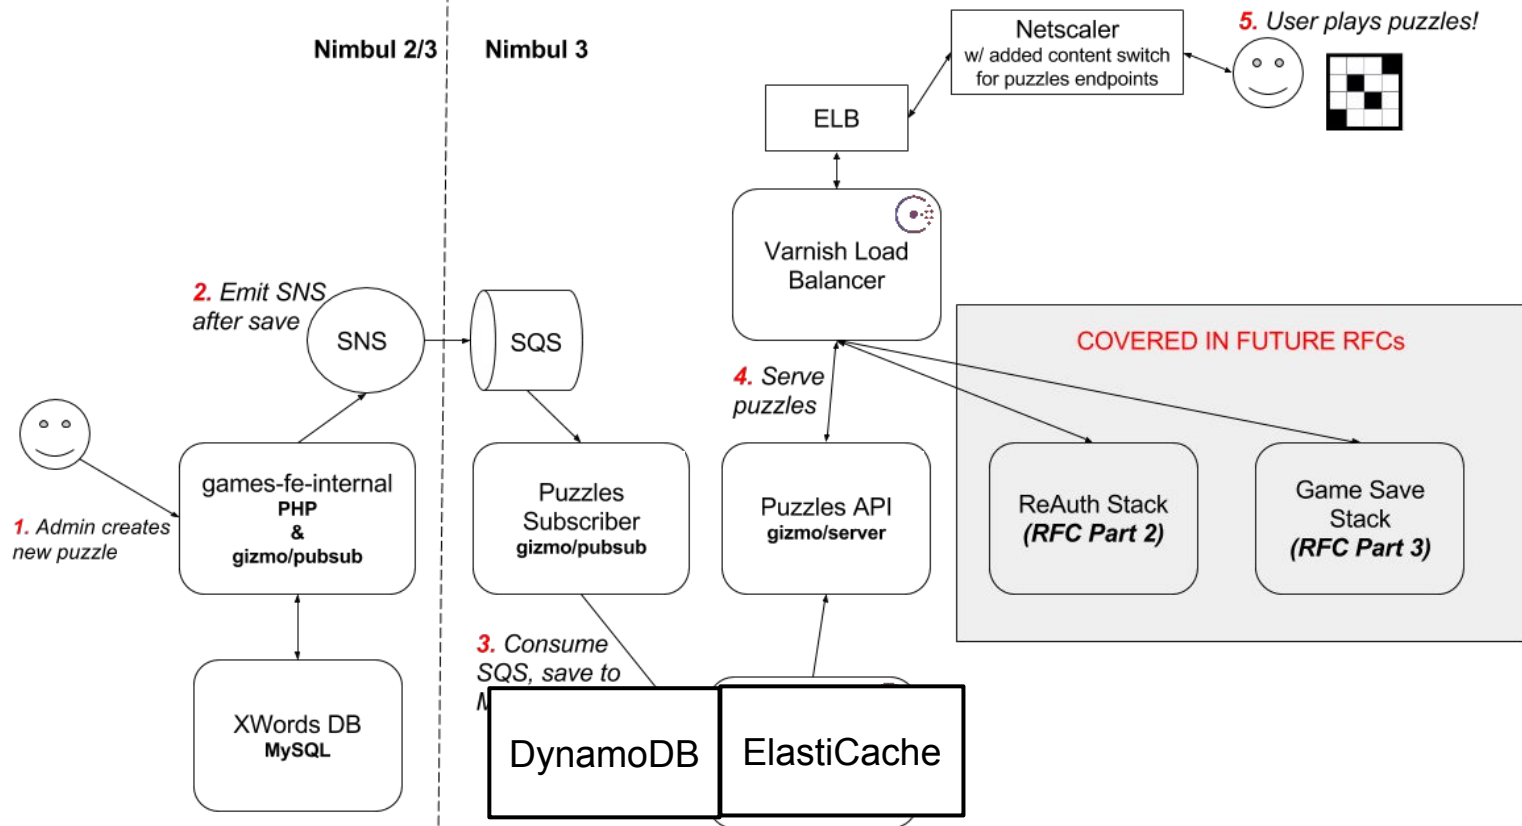
Deciding on Microservices

- Wanted to avoid any possible SPOF
 - Whatever happens, *always* serve the puzzle
- Needed to prototype and move quickly
 - Try new technology and swap out implementations
- Needed a temporary polyglot system
 - We can't rewrite it all at once!

Initial Architecture in AWS



Initial Architecture in AWS



Problems with AWS (for NYT)

- Alerting & Observability
- Service Discovery
- 0 downtime deployments
- Local development

New CTO, New Cloud Vision

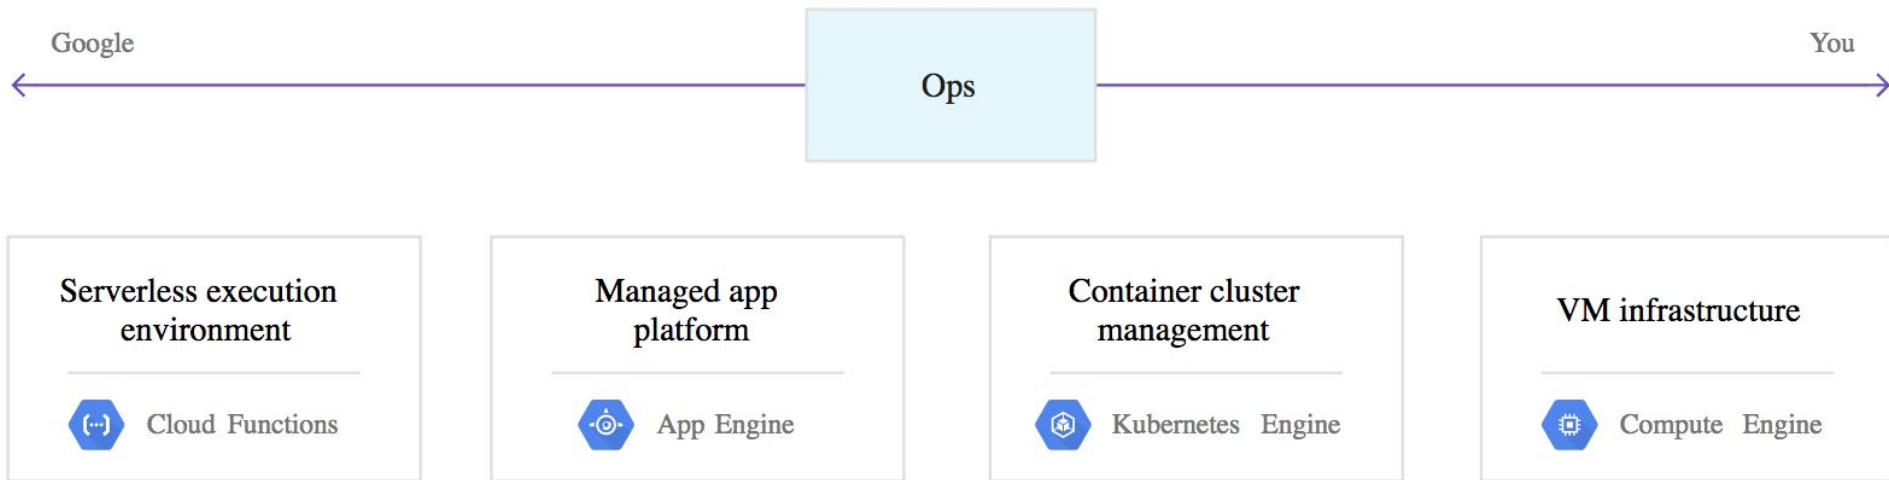
- Nick Rockwell joined Nov. 2015
- Moved everything to the cloud
- Introduced Fastly, saved 💰
- Expressed interest in Google Cloud Platform
- ***“Serverless”***

Serverless, According to Nick

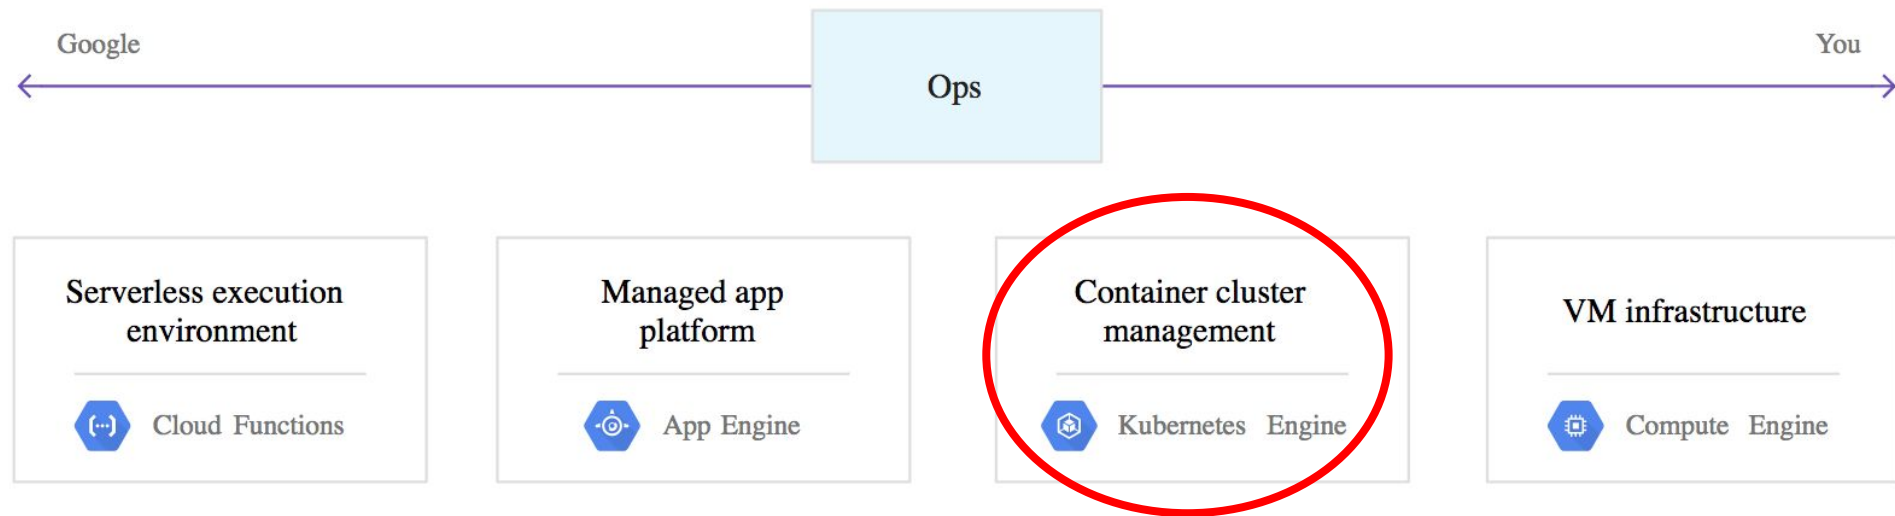
...serverless for me means any platform that provides utility while completely abstracting scaling and reliability.

- Nick Rockwell, CTO - NYTimes

Google's Compute Spectrum



Google's Compute Spectrum



Google Container Engine

scratch images, cool!

```
FROM scratch
```

```
ADD ca-certificates.crt /etc/ssl/certs/ca-certificates.crt
```

```
ADD zoneinfo.zip /usr/local/go/lib/time/zoneinfo.zip
```

```
ADD server /server
```

```
ENTRYPOINT ["/server"]
```

Google Container Registry

GCR is nice and simple!

```
gcloud docker [--authorize-only, -a] [--docker-host=DOCKER_HOST]  
  [--server=SERVER, [SERVER, ...], -s SERVER, [SERVER, ...];  
  default="gcr.io,us.gcr.io,eu.gcr.io,asia.gcr.io,b.gcr.io,  
  bucket.gcr.io,l.gcr.io,launcher.gcr.io,appengine.gcr.io,  
  us-mirror.gcr.io,eu-mirror.gcr.io,asia-mirror.gcr.io,mirror.gcr.  
  [GLOUD_WIDE_FLAG ...] [-- DOCKER_ARGS ...]
```

Google Container Engine



Google Container Engine



See a few
Kelsey Hightower
talks...

Google Container Engine!

OK, we have a service!!

...but something's missing

Google Container Engine

How do we monitor services?

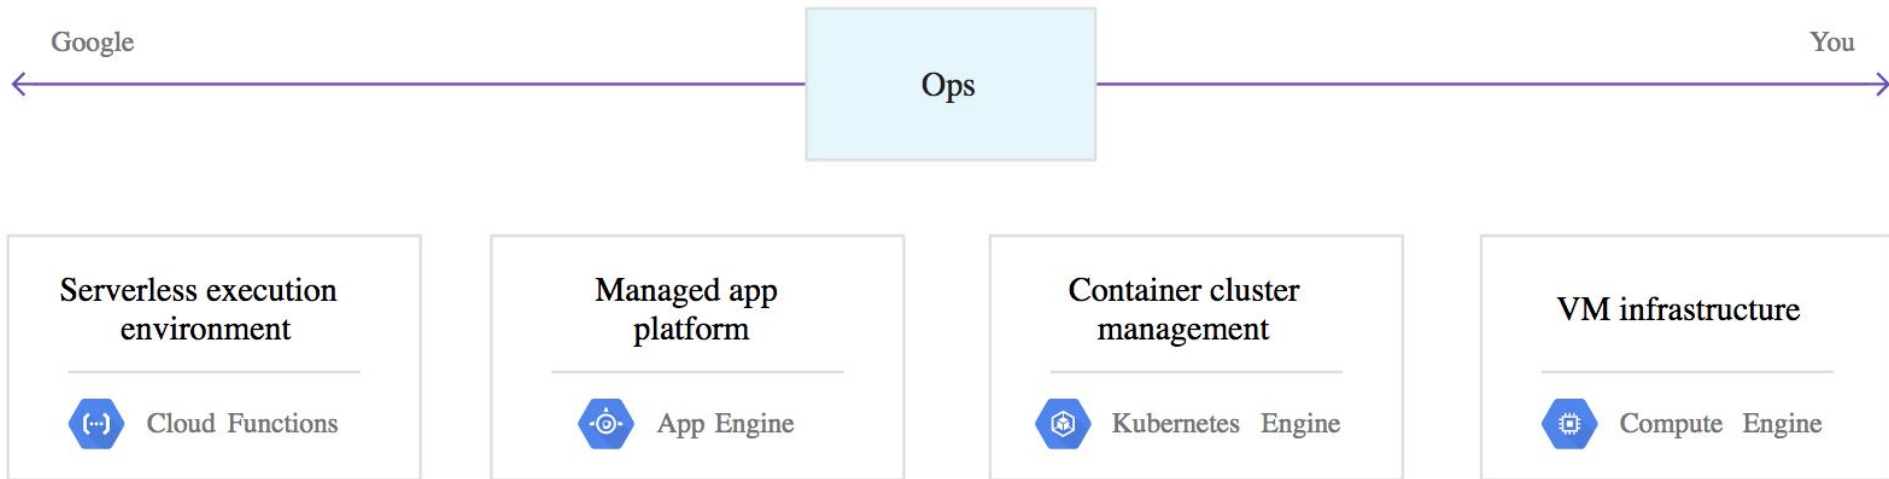


Google Container Engine?

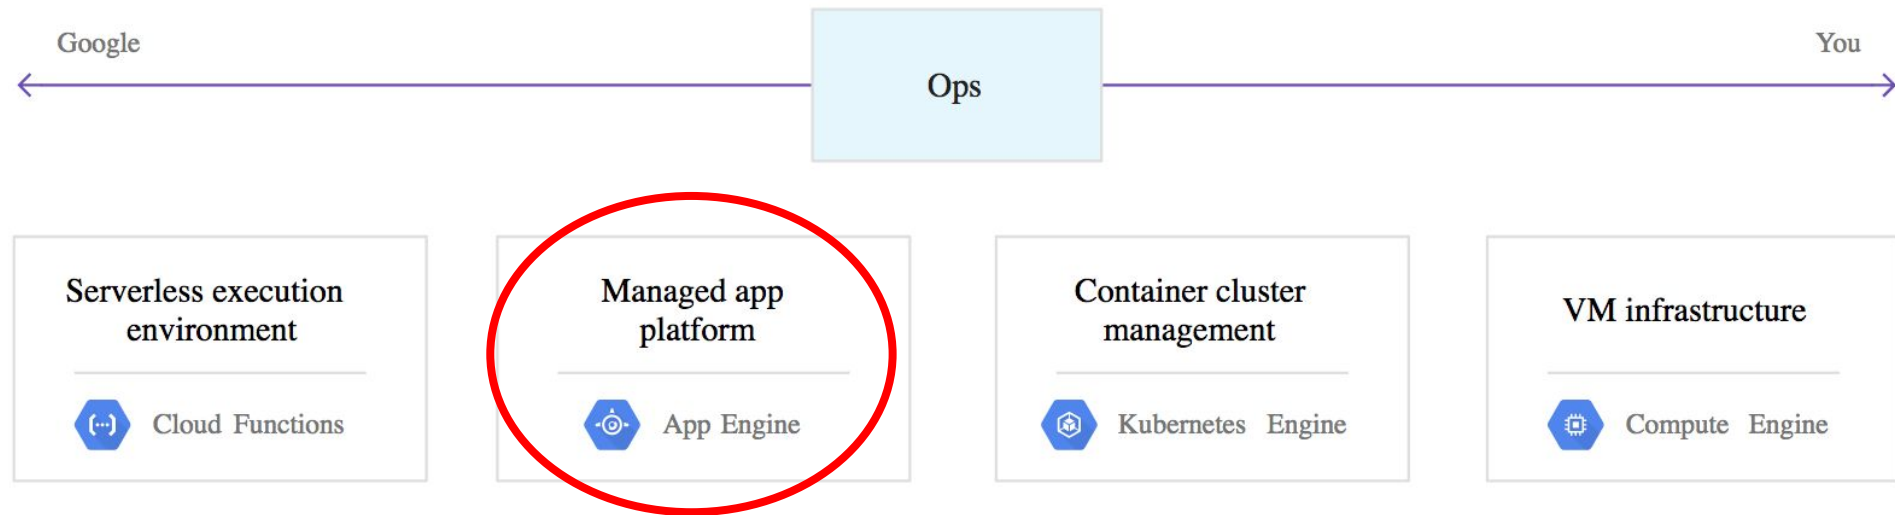
We have no way of automating SSL??

Bummer.

Google's Compute Spectrum



Google's Compute Spectrum



App Engine “Managed VMs”

I can reuse my container
...or it builds it for me!

```
└─ gcloud app deploy -h
Usage: gcloud app deploy [DEPLOYABLES ...] [optional flags]
optional flags may be  --bucket | --help | --image-url | --promote |
                        --stop-previous-version | --version | -h
```

App Engine “Managed VMs”

DNS and SSL for free!!

`https://SERVICE_ID-dot-MY_PROJECT_ID.appspot.com`

App Engine “Managed VMs”

It autoscales!

```
automatic_scaling:  
  min_num_instances: 1  
  max_num_instances: 15  
  cool_down_period_sec: 180  
  cpu_utilization:  
    target_utilization: 0.6
```


App Engine “Managed VMs”

It's not GA?! (2016)

...OK, let's revisit it later.

(note: it is GA *now*)

Google's Compute Spectrum



App Engine Standard

DNS and SSL!


`https://SERVICE_ID-dot-MY_PROJECT_ID.appspot.com`

App Engine Standard

Combined Access & App Logs!!


```
▼ 13:48:40.346 PUT 200 268 B 114 ms Firefox 52 /svc/crosswords/v2/game.json?regi_
0.1.0.40 - - [23/Jun/2017:13:48:40 -0400] "PUT /svc/crosswords/v2/game.json?regi_id=56254915
6/25 "Mozilla/5.0 (Windows NT 6.1; WOW64; rv:52.0) Gecko/20100101 Firefox/52.0 AppEngine-Goo
prd.appspot.com" ms=114 cpu_ms=46 cpm_usd=2.9951e-8 loading_request=0 instance=00c61b117cc28
8 trace_id=-
▼ {
  ▶ protoPayload: {...}
    insertId: "594d54790002b2e671eedcc0"
  ▶ httpRequest: {...}
  ▶ resource: {...}
    timestamp: "2017-06-23T17:48:40.346263Z"
    severity: "DEBUG"
  ▶ labels: {...}
    logName: "projects/nyt-games-prd/logs/appengine.googleapis.com%2Frequest_log"
  ▶ operation: {...}
    receiveTimestamp: "2017-06-23T17:48:41.178909151Z"
}
13:48:40.350 incoming commit: n7aroa
13:48:40.350 incoming update_timestamp: 1498239820
13:48:40.350 looking up progress for 51945265-13423
13:48:40.385 querying for newer commits that are also reset with key of 51945265-13423
13:48:40.390 skipping getall and applying to existing progress
13:48:40.395 adding commit: n7aroa
```

App Engine Standard Tracing!



App Engine Standard

Scales to 100s of instances in seconds?!



App Engine Standard

Simple user and service level security built in??

```
handlers:  
- url: /youraccount/.  
  script: _go_app  
  login: required  
  secure: always
```

App Engine Standard

Asserting identity to other App Engine apps



If you want to determine the identity of the App Engine app that is making a request to your App Engine app, you can use the request header X-Appengine-Inbound-Appid. This header is added to the request by the URLFetch service and is not user modifiable, so it safely indicates the requesting application's ID, if present.

App Engine Standard

Local Development

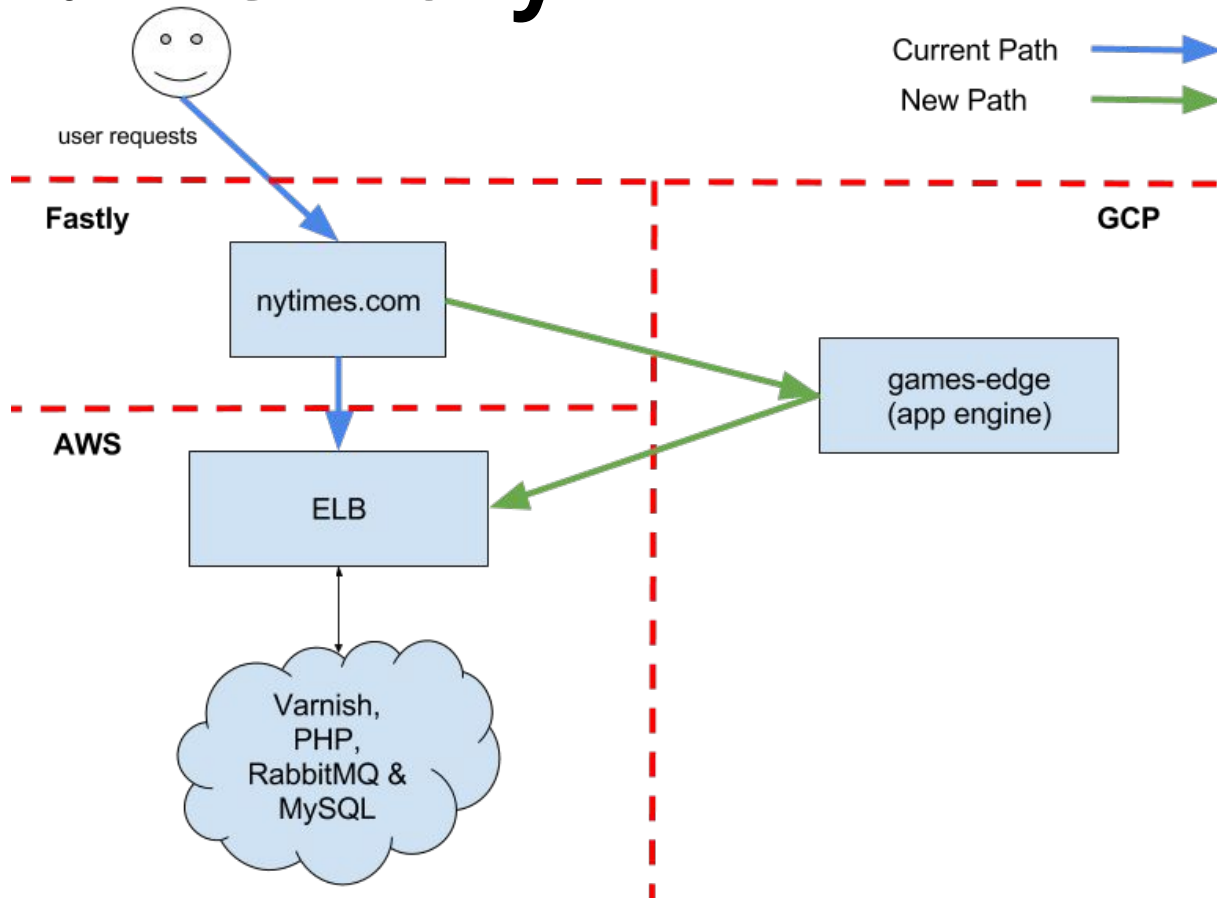
```
└─ goapp serve -h  
usage: serve [serve flags] [application_dir | package | yaml_files...]  
  
Serve launches your application on a local development App Engine server.
```

Time to Propose a GCP Solution

We wrote an internal RFC:

- 4 step process
- Most concerns were over lock in
- SLOs were made
- Agreed this will be a PoC

Step 1: The Proxy



Step 1: Problems

- Large headers
 - Too many third party cookies from web
 - Fix: Only pass what we need
- Autoscaling couldn't *quite* keep up
 - *"Request was aborted after waiting too long to attempt to service your request."*
 - Fix: sprinkle a little cron on it

Scheduled Autoscaling

- `description`: nightly 10pm scale up event
`url`: `/_ah/push-handlers/scale/up`
`schedule`: every day 21:55
`timezone`: America/New_York
`target`: default
- `description`: nightly scale down after 10pm scale up event
`url`: `/_ah/push-handlers/scale/down`
`schedule`: every day 22:05
`timezone`: America/New_York
`target`: default

Scheduled Autoscaling

```
- url: /_ah/push-handlers/*  
  script: _go_app  
  login: admin
```

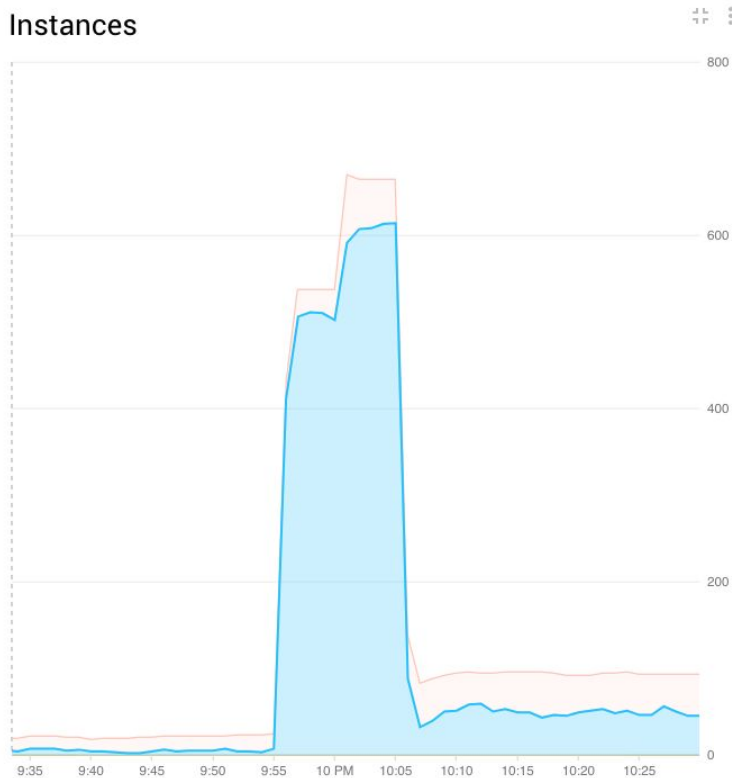
```
env_variables:  
  IDLE_INSTANCES_UP: 400  
  IDLE_INSTANCES_DOWN: 2
```

Scheduled Autoscaling

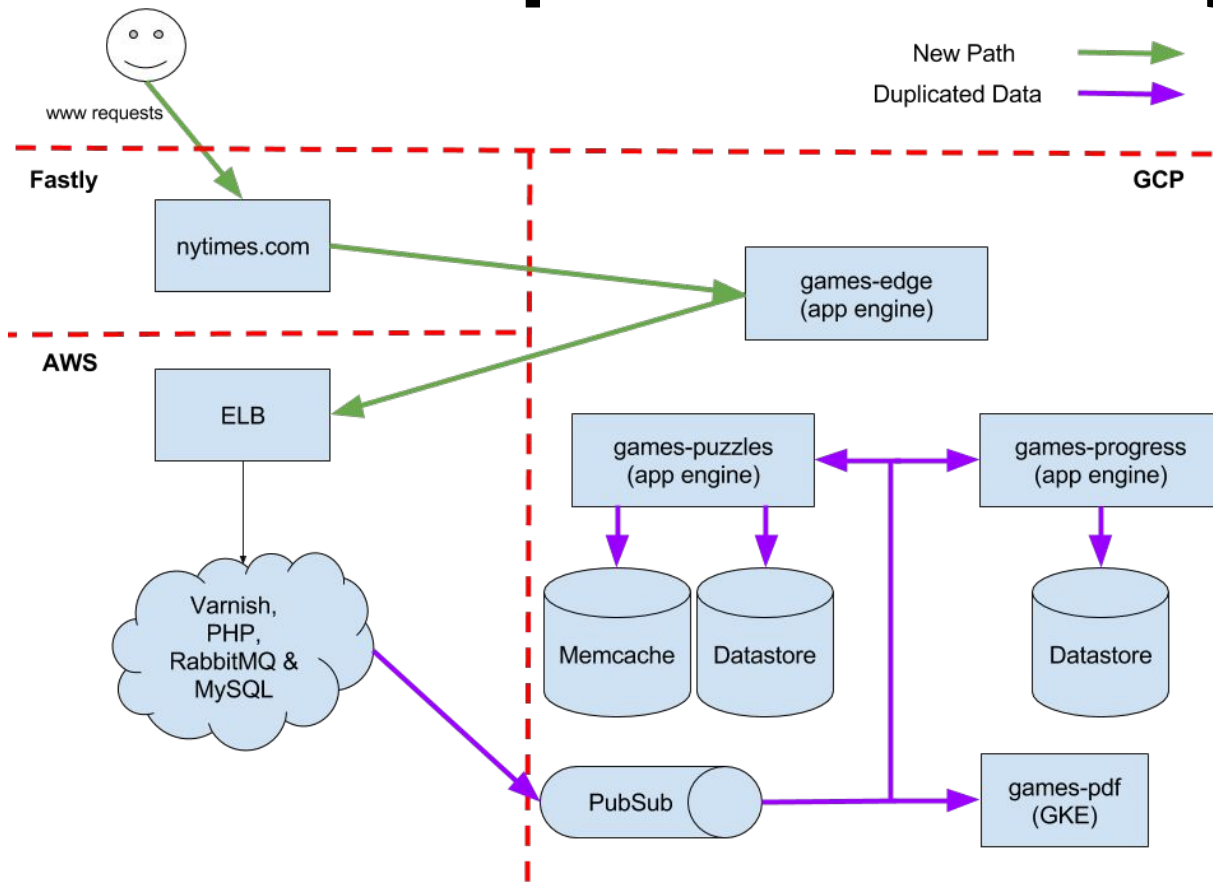
```
svc := appengine.NewAppsServicesVersionsService(clnt)
call := svc.Patch(appengin.AppID(ctx), svcName, vrsn, &appengine.Version{
    AutomaticScaling: &appengine.AutomaticScaling{
        MaxIdleInstances: idleInstances + 5,
        MinIdleInstances: idleInstances,
    },
}).UpdateMask("automaticScaling.min_idle_instances,automaticScaling.max_idle_instances")
```

Scheduled Autoscaling

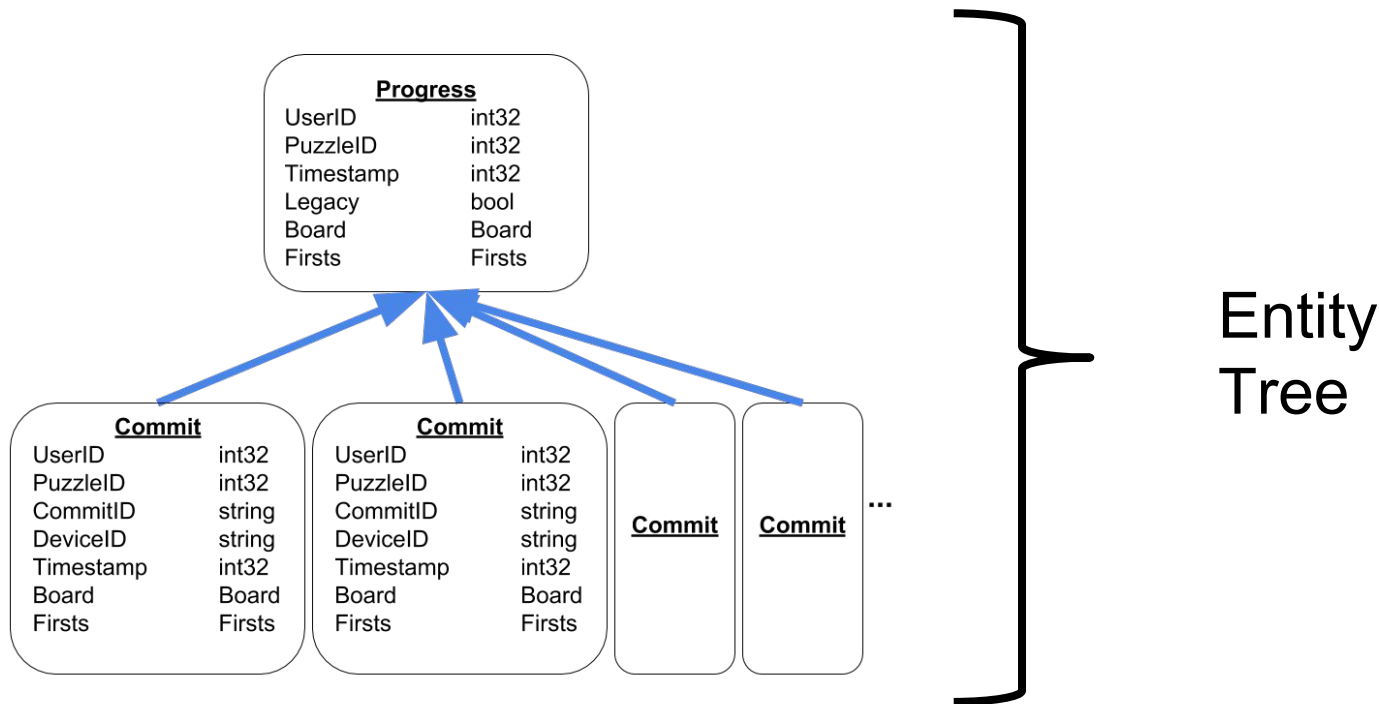
100s of instances in seconds!!!



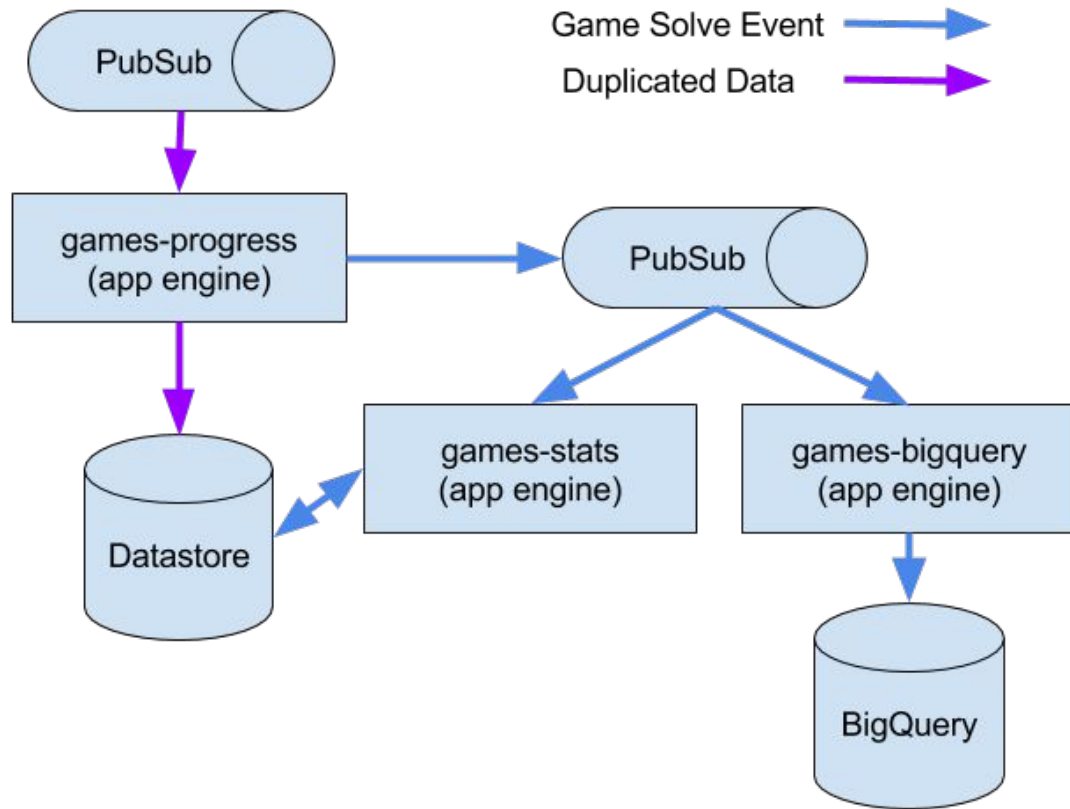
Step 2: New Endpoints + Data Sync



Step 2: Using Datastore



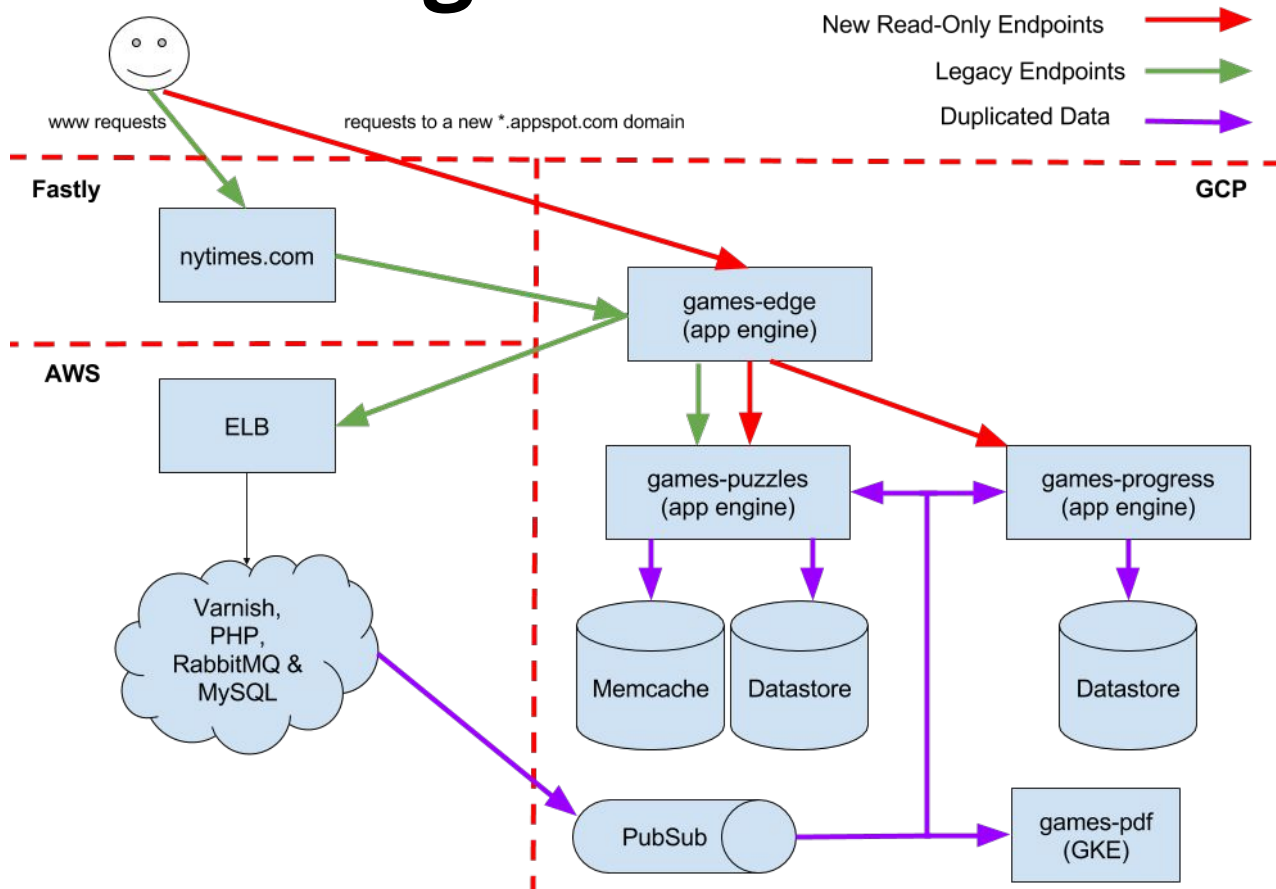
Step 2: Data Sync



Step 2: Problems

- How to shrink a PDF?
 - We need ghostscript!
 - Fix: GKE and PubSub
- Whoa! Datastore got pricey!
 - Too many reads when replaying data
 - Fix: add a lock around stats calculations

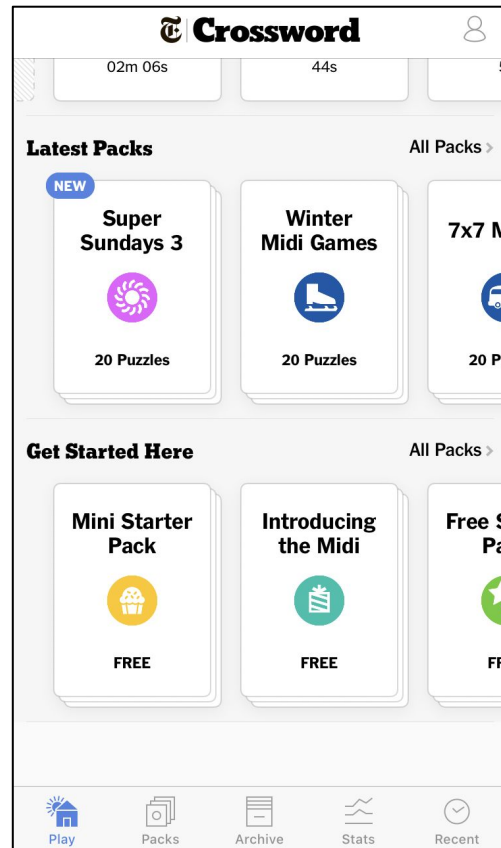
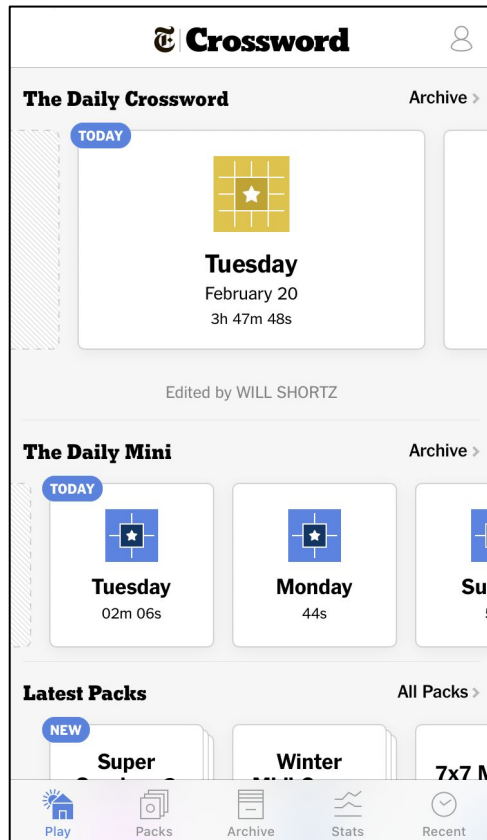
Step 3: Making GCP Live!



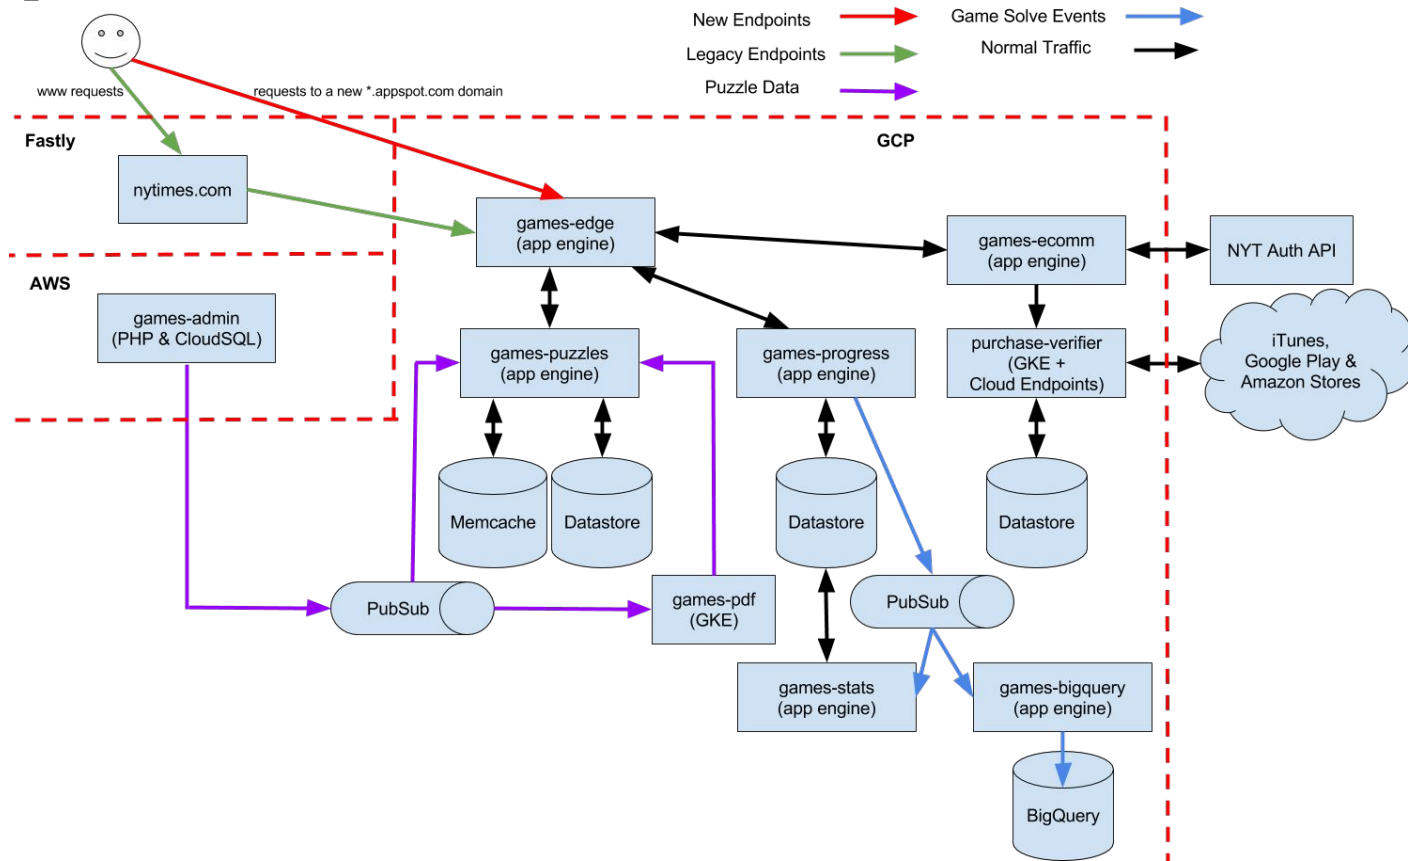


Step 3: The “Server View Model”

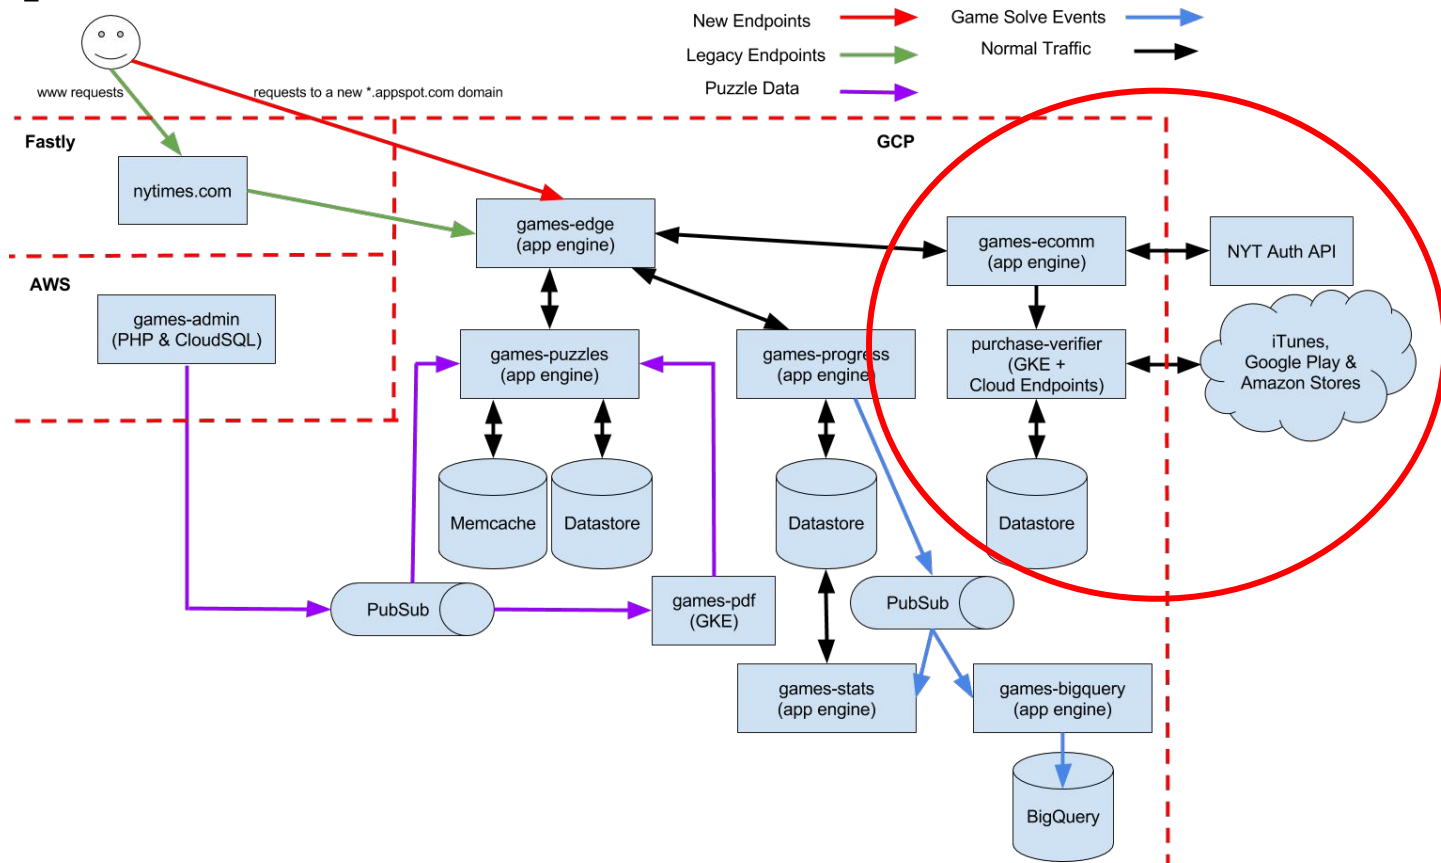
New iOS
home screen!



Step 4: Final Piece of the Puzzle



Step 4: Final Piece of the Puzzle



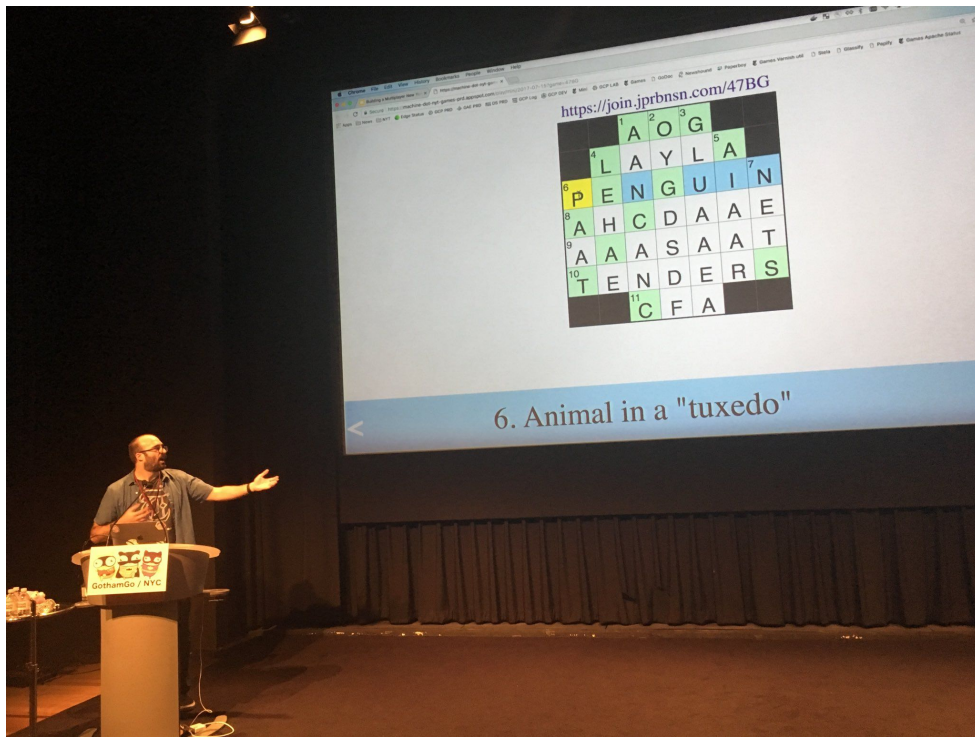
Step 4: OpenAPI, gRPC & Endpoints

- Service specs written in Open API
- Converted to Protobuf w/ openapi2proto
- Protos uploaded to Cloud Endpoints
- CE converts JSON to gRPC for us

The Results

- Started December 2016
 - ~10MM requests/day
- Completed September 2017
 - ~30MM+ requests/day
- Cloud costs cut in half
 - 💰💰💰💰
- 6 game prototypes publicly tested
 - Coming soon: Spelling Bee!

Multiplayer Crosswords!



What's Next?



JP Robinson

@jprbnsn

been playing around with @googlecloud's #spanner, #appengine and #golang to manage (new!) social leaderboards for the @NYT_Crossword this week.

great dev experience so far, but I should probably reach out to our design team 🤔

mini - 2017-12-22

Name	Timer	# Revealed
JP	22	0
Ben	26	0

10:54 AM - 22 Dec 2017

Thanks!

Questions?